


1. Executive Summary

The deployment of Generative AI in the financial sector faces a critical barrier: the non-deterministic nature of Large Language Models (LLMs) fundamentally conflicts with the zero-tolerance compliance requirements of financial services. This whitepaper proposes a generic 5-Layer Reference Architecture designed to decouple generative AI capabilities from strict regulatory constraints. By shifting governance from pure algorithmic training to a structural interaction framework, this architecture provides Small and Medium-Sized Enterprises (SMEs) with a plug-and-play blueprint to achieve human-AI symbiosis without sacrificing regulatory compliance.

2. The 5-Layer Trust & Routing Architecture



Layer 1: The Presentation & Identity Layer (L1)

Establishes a unified, single-persona agent identity. Rather than presenting as a generic bot, the agent maintains a consistent visual and interaction baseline, ensuring users have a predictable focal point for trust during high-stakes financial inquiries.

Layer 2: The Affective & Tone Configuration Layer (L2)

Replaces rigid robotic responses with an engineered emotional categorization system. This layer dynamically adjusts the agent's tone based on the user's scenario matrix (e.g., New User Onboarding vs. Overdue Bill Inquiry). It maps conversational inputs to specific affective states:

Guide & Delight: For standard navigational and positive reinforcement scenarios.

Remind (Gentle Urgent): For early-stage delinquency or risk warnings.

Support (Reassuring & Crisis): For account freezing, disputes, or escalation events, triggering immediate visual and tonal shifts to stabilize user sentiment.

Layer 3: The Functional Role-Switching Layer (L3)

Implements a dual-mode "Temperature Control" mechanism to separate roles:

Mode 0 (Strict Compliance/CS): Deterministic Standard Operating Procedures (SOPs) for resolving hard constraints (e.g., refunds, legal disputes).

Mode 1 (Marketing/Engagement): Free-form conversational marketing for product recommendations and user engagement.

The system allows seamless switching between these modes within the same session lifecycle based on intent recognition.

Layer 4: The Decision & Boundary Routing Layer (L4)

The core "insurance fuse" of the architecture. It defines strict boundaries between AI-made and Human-made decisions:

AI-Driven Free Zone: Handles soft SOPs and general inquiries with dynamic LLM generation.

Escalation & Fallback: When a user's intent hits a predefined risk matrix (e.g., severe complaint, fraud suspicion), the routing layer instantly revokes the LLM's generative privileges. It forces a fallback to deterministic rule engines or smoothly hands off the session to a Human-in-the-Loop (HITL) without causing UX friction.

Layer 5: The Evaluation & Audit Trail Layer (L5)

A continuous monitoring framework ensuring the "black box" remains transparent to regulators. It evaluates AI performance across three dimensions:

Response Content: Real-time checking of relevance and factual accuracy against financial databases.

Trigger Consistency: Ensuring that emotional shifts and SOP invocations match the pre-defined scenario matrices perfectly.

Auditability: Generating immutable, CFPB-compliant logs for every automated decision and fallback trigger.